

# ECE763 Computer Vision: Models, Learning and Inference (including Deep Learning)



## Project 03

**Due (see moodle) No late days**  
(Grades: 100 points in total)

Depart. of ECE, NC State University

Instructor: Tianfu (Matt) Wu



Electrical and  
Computer Engineering

0



## Grading Policy

- **Homework:** 7% x 5 = 35%
- **Projects:** 60% + [5% bonus]
  - Project 1 (10%) & 2 (20%): **work individually [not required if you work on thrust 1 and 2 in the next slide]**
  - Project 3 (course project): 30%, **a group (<= 3 members) is allowed [not required for all thrusts in the next slide]**
    - You can propose your own projects and work on it once agreed by both of you and me.
    - Requirement: final write-up (15%) and self-contained reproducible code (15%)
    - Bonus points (5%): for top-3 winner of the final project
- **Attendance:** 5% (via in-class moodle quiz, 10min before and 15min after class, from the 2<sup>nd</sup> week)
- **Late policy**
  - 5 free late days – use them in your ways (counted using 0.5 unit, <=6 hours as 0.5 late day, otherwise 1 later day)
  - Afterwards, 25% off per day late
  - Not accepted after 3 late days per HW and Project 1 & 2
  - Does not apply to Project 3 (no late assignments allowed).
  - **We will NOT accept any replacement of submission after deadline, even if you can show the time stamp of the replacement is earlier than the deadline. So, please double-check if you submit correct files.**
- **Other policies:** please see details in the syllabus.

1

1



## Discussions on Projects

- Thrust I: If applicable, **identify problem(s)** in your own field that have images/videos generated and to be analyzed, and **collect dataset(s)** (e.g., publicly available ones, or from your own research)
  - Formulate the problem(s): e.g., input-output mapping, target platforms (edge device or workstation/server), evaluation metrics
  - Find a few references in top-tier conferences/journals
  - Our goal: to have a submission ready by the end of this class, through customized project 1-3 design
- Thrust II: If no specific domains of interest,
  - Check <https://paperswithcode.com/>
  - Find problem(s) that interest you
  - Survey SOTA methods therein
  - Our goal: to have a submission to top-tier CV conference ready by the end of this class, through customized project 1-3 design
- Thrust III: If both do not work for you
  - We will have standard project assignments.

2

2



## Instructions and Notes

- *How to submit your solutions:* put your report (word or pdf) and results images (.png) if had in a folder named [your\_unityid]\_hw01 (e.g., twu19\_hw01), and then compress it as a zip file (e.g., twu19\_hw01.zip). Submit the zip file through **moodle**.
- *If you miss the deadline and still have unused late days, please send your zip file to TAs and me.*
  - 5 free late days (counted using 0.5 unit,  $\leq 6$  hours as 0.5 late day, otherwise 1 later day) in total – use them in your ways; Afterwards, 25% off per day late; **Not** accepted after 3 late days per HW and Project. Not applicable to the final project.
- **Important Note:** We will **NOT** accept any replacement of submission after deadline ([+late days you use]), even if you can show the time stamp of the replacement is earlier than the deadline. So, **please double-check whether you submit correct files.**

3

3



## Project 03

- Project 03: 30%
  - Requirement:
    - final write-up (15%) and
    - self-contained reproducible code (15%): **if your code is not reproducible, you will lose 10%; So, please make sure what you report in the write-up are consistent with the code.**
  - Bonus points: 5
    - For novel ideas (including novel implementation of some existing techniques, e.g., in a significantly faster way under fair comparisons).
    - If you want to claim this, please add a **section 0** in your write-up to clearly present and justify **what's new**. **The novelty will be checked and judged by both TAs and the instructor.**

4

4



## Babysitting the training of DNN

- **Problem:** Practice the babysitting method of training neural network as discussed in Lecture notes 21-24.
  - **Step 0:** Data (you can start with existing data provided in different DL code platform, e.g., MNIST in PyTorch), and **you must test your code using the face data you used in project 1 or 2.**
  - **Step 1:** preprocess the data and select a data augment scheme. (Note: you can compare w/ and w/o this preprocessing step and the data-aug to see how they affect your training and testing performance)
  - **Step 2:** Choose the architecture. E.g. you can use something simple, a 3-layer FC network or the LeNet 5 (available in all deep learning platforms).
  - > Then please follow the slides (see the next slides 6 - 25). to output detailed information of this babysitting procedure.
- **Hint:** You can reuse the tutorial code in different deep learning platform (e.g., the MNIST / CIFAR10 tutorial is available in almost all platforms).
  - E.g., If you use pytorch, you can find a lot of examples at <https://pytorch.org/tutorials/>
- **Requirement:** Although you can reuse the tutorial code, you need to modify the code to output different babysitting information (see the next slides 4 - 23). You need to snap your screenshots and paste those in your reports as shown in the slides. You also need to provide **self-contained reproducible code**.

5

5

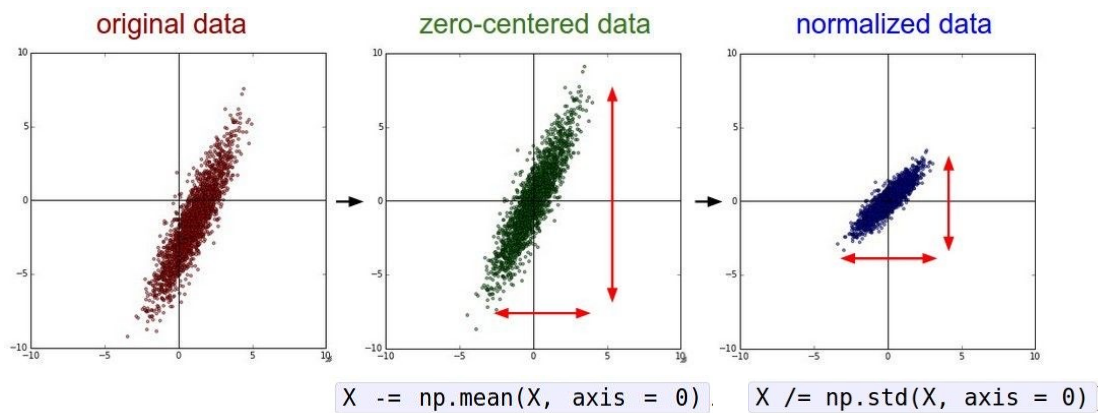


## Project 03 -- Babysitting the Learning Process

6



## Step 1 Data Preprocessing



(Assume  $X$  [NxD] is data matrix,  
each example in a row)

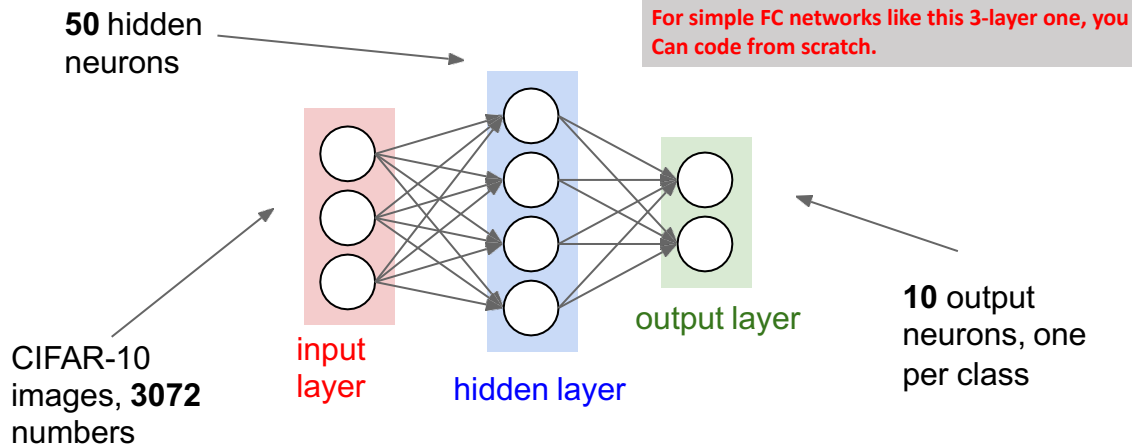
You can try: with and w/o data normalization

7



## Step 2: Choose the architecture

- E.g., we start with one hidden layer of 50 neurons or LeNet or AlexNet ...



8



## Double check that the loss is reasonable:

CHECK this step with your network and data

```
def init_two_layer_model(input_size, hidden_size, output_size):
    # initialize a model
    model = {}
    model['W1'] = 0.0001 * np.random.randn(input_size, hidden_size)
    model['b1'] = np.zeros(hidden_size)
    model['W2'] = 0.0001 * np.random.randn(hidden_size, output_size)
    model['b2'] = np.zeros(output_size)
    return model
```

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
loss, grad = two_layer_net(X_train, model, y_train, 0.0)
print loss
```

2.30261216167

loss ~2.3.  
"correct" for  
10 classes

returns the loss and the  
gradient for all parameters

disable regularization

9



## Double check that the loss is reasonable:

```
def init_two_layer_model(input_size, hidden_size, output_size):
    # initialize a model
    model = {}
    model['W1'] = 0.0001 * np.random.randn(input_size, hidden_size)
    model['b1'] = np.zeros(hidden_size)
    model['W2'] = 0.0001 * np.random.randn(hidden_size, output_size)
    model['b2'] = np.zeros(output_size)
    return model
```

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
loss, grad = two_layer_net(X_train, model, y_train, 1e3)
print loss
```

crank up regularization

CHECK this step with your network and data

3.06859716482

loss went up, good. (sanity check)

10



## Lets try to train now...

**Tip:** Make sure that you can overfit very small portion of the training data

CHECK this step with your network and data

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
X_tiny = X_train[:20] # take 20 examples
y_tiny = y_train[:20]
best_model, stats = trainer.train(X_tiny, y_tiny, X_tiny, y_tiny,
                                  model, two_layer_net,
                                  num_epochs=200, reg=0.0,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = False,
                                  learning_rate=1e-3, verbose=True)
```

The above code:

- take the first 20 examples from CIFAR-10
- turn off regularization (reg = 0.0)
- use simple vanilla 'sgd'

11





## Lets try to train now...

**Tip:** Make sure that you can overfit very small portion of the training data

Very small loss,  
train accuracy 1.00,  
nice!

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
X_tiny = X_train[:20] # take 20 examples
y_tiny = y_train[:20]
best_model, stats = trainer.train(X_tiny, y_tiny, X_tiny, y_tiny,
                                  model, two_layer_net,
                                  num_epochs=200, reg=0.0,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = False,
                                  learning_rate=1e-3, verbose=True)
```

```
Finished epoch 1 / 200: cost 2.302603, train: 0.400000, val 0.400000, lr 1.000000e-03
Finished epoch 2 / 200: cost 2.302258, train: 0.450000, val 0.450000, lr 1.000000e-03
Finished epoch 3 / 200: cost 2.301849, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 4 / 200: cost 2.301196, train: 0.650000, val 0.650000, lr 1.000000e-03
Finished epoch 5 / 200: cost 2.300044, train: 0.650000, val 0.650000, lr 1.000000e-03
Finished epoch 6 / 200: cost 2.297864, train: 0.550000, val 0.550000, lr 1.000000e-03
Finished epoch 7 / 200: cost 2.293595, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 8 / 200: cost 2.285096, train: 0.550000, val 0.550000, lr 1.000000e-03
Finished epoch 9 / 200: cost 2.268094, train: 0.550000, val 0.550000, lr 1.000000e-03
Finished epoch 10 / 200: cost 2.234787, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 11 / 200: cost 2.173187, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 12 / 200: cost 2.076862, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 13 / 200: cost 1.974090, train: 0.400000, val 0.400000, lr 1.000000e-03
Finished epoch 14 / 200: cost 1.895885, train: 0.400000, val 0.400000, lr 1.000000e-03
Finished epoch 15 / 200: cost 1.820876, train: 0.450000, val 0.450000, lr 1.000000e-03
Finished epoch 16 / 200: cost 1.737430, train: 0.450000, val 0.450000, lr 1.000000e-03
Finished epoch 17 / 200: cost 1.642356, train: 0.500000, val 0.500000, lr 1.000000e-03
Finished epoch 18 / 200: cost 1.535239, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 19 / 200: cost 1.421527, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 20 / 200: cost 1.265750, train: 0.600000, val 0.600000, lr 1.000000e-03
Finished epoch 195 / 200: cost 0.002694, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 196 / 200: cost 0.002674, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 197 / 200: cost 0.002655, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 198 / 200: cost 0.002635, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 199 / 200: cost 0.002617, train: 1.000000, val 1.000000, lr 1.000000e-03
Finished epoch 200 / 200: cost 0.002597, train: 1.000000, val 1.000000, lr 1.000000e-03
finished optimization. best validation accuracy: 1.000000
```

12



## Lets try to train now...

Start with small regularization and find learning rate that makes the loss go down.

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=1e-6, verbose=True)
```

```
Finished epoch 1 / 10: cost 2.302576, train: 0.080000, val 0.103000, lr 1.000000e-06
Finished epoch 2 / 10: cost 2.302582, train: 0.121000, val 0.124000, lr 1.000000e-06
Finished epoch 3 / 10: cost 2.302558, train: 0.119000, val 0.138000, lr 1.000000e-06
Finished epoch 4 / 10: cost 2.302519, train: 0.127000, val 0.151000, lr 1.000000e-06
Finished epoch 5 / 10: cost 2.302517, train: 0.158000, val 0.171000, lr 1.000000e-06
Finished epoch 6 / 10: cost 2.302518, train: 0.179000, val 0.172000, lr 1.000000e-06
Finished epoch 7 / 10: cost 2.302466, train: 0.180000, val 0.176000, lr 1.000000e-06
Finished epoch 8 / 10: cost 2.302452, train: 0.175000, val 0.185000, lr 1.000000e-06
Finished epoch 9 / 10: cost 2.302459, train: 0.206000, val 0.192000, lr 1.000000e-06
Finished epoch 10 / 10: cost 2.302420, train: 0.190000, val 0.192000, lr 1.000000e-06
finished optimization. best validation accuracy: 0.192000
```

CHECK this step with your network and data

13



## Lets try to train now...

Start with small regularization and find learning rate that makes the loss go down.

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches=True,
                                  learning_rate=1e-6, verbose=True)
```

Epoch	Cost	Train Loss	Val Loss	Learning Rate
Finished epoch 1 / 10:	2.302576	0.080000	0.103000	1.000000e-06
Finished epoch 2 / 10:	2.302582	0.121000	0.124000	1.000000e-06
Finished epoch 3 / 10:	2.302558	0.119000	0.138000	1.000000e-06
Finished epoch 4 / 10:	2.302519	0.127000	0.151000	1.000000e-06
Finished epoch 5 / 10:	2.302517	0.158000	0.171000	1.000000e-06
Finished epoch 6 / 10:	2.302518	0.179000	0.172000	1.000000e-06
Finished epoch 7 / 10:	2.302466	0.180000	0.176000	1.000000e-06
Finished epoch 8 / 10:	2.302452	0.175000	0.185000	1.000000e-06
Finished epoch 9 / 10:	2.302459	0.206000	0.192000	1.000000e-06
Finished epoch 10 / 10:	2.302420	0.190000	0.192000	1.000000e-06

finished optimization. best validation accuracy: 0.192000

Loss barely changing

14



## Lets try to train now...

Start with small regularization and find learning rate that makes the loss go down.

loss not going down:  
learning rate too low

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches=True,
                                  learning_rate=1e-6, verbose=True)
```

Epoch	Cost	Train Loss	Val Loss	Learning Rate
Finished epoch 1 / 10:	2.302576	0.080000	0.103000	1.000000e-06
Finished epoch 2 / 10:	2.302582	0.121000	0.124000	1.000000e-06
Finished epoch 3 / 10:	2.302558	0.119000	0.138000	1.000000e-06
Finished epoch 4 / 10:	2.302519	0.127000	0.151000	1.000000e-06
Finished epoch 5 / 10:	2.302517	0.158000	0.171000	1.000000e-06
Finished epoch 6 / 10:	2.302518	0.179000	0.172000	1.000000e-06
Finished epoch 7 / 10:	2.302466	0.180000	0.176000	1.000000e-06
Finished epoch 8 / 10:	2.302452	0.175000	0.185000	1.000000e-06
Finished epoch 9 / 10:	2.302459	0.206000	0.192000	1.000000e-06
Finished epoch 10 / 10:	2.302420	0.190000	0.192000	1.000000e-06

finished optimization. best validation accuracy: 0.192000

Loss barely changing: Learning rate is  
probably too low

15





## Lets try to train now...

Start with small regularization and find learning rate that makes the loss go down.

loss not going down:  
learning rate too low

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=1e-6, verbose=True)
```

Epoch	Cost	Train Accuracy	Val Accuracy	Learning Rate
Finished epoch 1 / 10:	2.302576	0.080000	0.103000	1.000000e-06
Finished epoch 2 / 10:	2.302582	0.121000	0.124000	1.000000e-06
Finished epoch 3 / 10:	2.302558	0.119000	0.138000	1.000000e-06
Finished epoch 4 / 10:	2.302519	0.127000	0.151000	1.000000e-06
Finished epoch 5 / 10:	2.302517	0.158000	0.171000	1.000000e-06
Finished epoch 6 / 10:	2.302518	0.179000	0.172000	1.000000e-06
Finished epoch 7 / 10:	2.302466	0.180000	0.176000	1.000000e-06
Finished epoch 8 / 10:	2.302452	0.175000	0.185000	1.000000e-06
Finished epoch 9 / 10:	2.302459	0.206000	0.192000	1.000000e-06
Finished epoch 10 / 10:	2.302420	0.190000	0.192000	1.000000e-06

finished optimization. best validation accuracy: 0.192000

Loss barely changing: Learning rate is probably too low

Notice train/val accuracy goes to 20% though, what's up with that? (remember this is softmax)

16



## Lets try to train now...

• Start with small regularization and find learning rate that makes the loss go down.

• loss not going down:

- learning rate too low

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
```

Now let's try learning rate 1e6.

April 19, 2018

17



## Lets try to train now...

Start with small regularization and find learning rate that makes the loss go down.

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=1e6, verbose=True)

/home/karpathy/cs231n/code/cs231n/classifiers/neural_net.py:50: RuntimeWarning: divide by zero encountered in log
  data_loss = -np.sum(np.log(probs[range(N), y])) / N
/home/karpathy/cs231n/code/cs231n/classifiers/neural_net.py:48: RuntimeWarning: invalid value encountered in subtract
  probs = np.exp(scores - np.max(scores, axis=1, keepdims=True))
Finished epoch 1 / 10: cost nan, train: 0.091000, val 0.087000, lr 1.000000e+06
Finished epoch 2 / 10: cost nan, train: 0.095000, val 0.087000, lr 1.000000e+06
Finished epoch 3 / 10: cost nan, train: 0.100000, val 0.087000, lr 1.000000e+06
```

**loss not going down:**  
learning rate too low  
**loss exploding:**  
learning rate too high

cost: NaN almost  
always means high  
learning rate...

18



## Lets try to train now...

Start with small regularization and find learning rate that makes the loss go down.

```
model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
trainer = ClassifierTrainer()
best_model, stats = trainer.train(X_train, y_train, X_val, y_val,
                                  model, two_layer_net,
                                  num_epochs=10, reg=0.000001,
                                  update='sgd', learning_rate_decay=1,
                                  sample_batches = True,
                                  learning_rate=3e-3, verbose=True)

Finished epoch 1 / 10: cost 2.186654, train: 0.308000, val 0.306000, lr 3.000000e-03
Finished epoch 2 / 10: cost 2.176230, train: 0.330000, val 0.350000, lr 3.000000e-03
Finished epoch 3 / 10: cost 1.942257, train: 0.376000, val 0.352000, lr 3.000000e-03
Finished epoch 4 / 10: cost 1.827868, train: 0.329000, val 0.310000, lr 3.000000e-03
Finished epoch 5 / 10: cost inf, train: 0.128000, val 0.128000, lr 3.000000e-03
Finished epoch 6 / 10: cost inf, train: 0.144000, val 0.147000, lr 3.000000e-03
```

**loss not going down:**  
learning rate too low  
**loss exploding:**  
learning rate too high

3e-3 is still too high. Cost explodes....

=> Rough range for learning rate we  
should be cross-validating is  
somewhere [1e-3 ... 1e-5]

19



# Hyperparameter Optimization

**CHECK this step (weight decay, learning rate) with your network and data, as well as other hyperparameters which you observe to be sensitive in your training.**

20



## Cross-validation strategy

**coarse -> fine** cross-validation in stages

**First stage:** only a few epochs to get rough idea of what params work

**Second stage:** longer running time, finer search

... (repeat as necessary)

Tip for detecting explosions in the solver:

If the cost is ever  $> 3 \times$  original cost, break out early

21



## For example: run coarse search for 5 epochs

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-5, 5)
    lr = 10**uniform(-3, -6)

    trainer = ClassifierTrainer()
    model = init_two_layer_model(32*32*3, 50, 10) # input size, hidden size, number of classes
    trainer = ClassifierTrainer()
    best_model_local, stats = trainer.train(X_train, y_train, X_val, y_val,
                                           model, two_layer_net,
                                           num_epochs=5, reg=reg,
                                           update='momentum', learning_rate_decay=0.9,
                                           sample_batches = True, batch_size = 100,
                                           learning_rate=lr, verbose=False)
```

note it's best to optimize in log space!

nice

```
val_acc: 0.412000, lr: 1.405206e-04, reg: 4.793564e-01, (1 / 100)
val_acc: 0.214000, lr: 7.231888e-06, reg: 2.321281e-04, (2 / 100)
val_acc: 0.208000, lr: 2.119571e-06, reg: 8.011857e+01, (3 / 100)
val_acc: 0.196000, lr: 1.551131e-05, reg: 4.374936e-05, (4 / 100)
val_acc: 0.079000, lr: 1.753300e-05, reg: 1.200424e+03, (5 / 100)
val_acc: 0.223000, lr: 4.215128e-05, reg: 4.196174e+01, (6 / 100)
val_acc: 0.441000, lr: 1.750259e-04, reg: 2.110807e-04, (7 / 100)
val_acc: 0.241000, lr: 6.749231e-05, reg: 4.226413e+01, (8 / 100)
val_acc: 0.482000, lr: 4.296863e-04, reg: 6.642555e-01, (9 / 100)
val_acc: 0.079000, lr: 5.401602e-06, reg: 1.599828e+04, (10 / 100)
val_acc: 0.154000, lr: 1.618508e-06, reg: 4.925252e-01, (11 / 100)
```

22



## Now run finer search...

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-5, 5)
    lr = 10**uniform(-3, -6)
```

adjust range

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-4, 0)
    lr = 10**uniform(-3, -4)
```

```
val_acc: 0.527000, lr: 5.340517e-04, reg: 4.097824e-01, (0 / 100)
val_acc: 0.492000, lr: 2.279484e-04, reg: 9.991345e-04, (1 / 100)
val_acc: 0.512000, lr: 8.680827e-04, reg: 1.349727e-02, (2 / 100)
val_acc: 0.461000, lr: 1.028377e-04, reg: 1.220193e-02, (3 / 100)
val_acc: 0.460000, lr: 1.113730e-04, reg: 5.244309e-02, (4 / 100)
val_acc: 0.498000, lr: 9.477776e-04, reg: 2.001293e-03, (5 / 100)
val_acc: 0.469000, lr: 1.484369e-04, reg: 4.328313e-01, (6 / 100)
val_acc: 0.522000, lr: 5.586261e-04, reg: 2.312685e-04, (7 / 100)
val_acc: 0.530000, lr: 5.808183e-04, reg: 8.259964e-02, (8 / 100)
val_acc: 0.489000, lr: 1.979168e-04, reg: 1.010889e-04, (9 / 100)
val_acc: 0.490000, lr: 2.036031e-04, reg: 2.406271e-03, (10 / 100)
val_acc: 0.475000, lr: 2.021162e-04, reg: 2.287807e-01, (11 / 100)
val_acc: 0.460000, lr: 1.135527e-04, reg: 3.905040e-02, (12 / 100)
val_acc: 0.515000, lr: 6.947668e-04, reg: 1.562808e-02, (13 / 100)
val_acc: 0.531000, lr: 9.471549e-04, reg: 1.433895e-03, (14 / 100)
val_acc: 0.509000, lr: 3.140888e-04, reg: 2.857518e-01, (15 / 100)
val_acc: 0.514000, lr: 6.438349e-04, reg: 3.033781e-01, (16 / 100)
val_acc: 0.502000, lr: 3.921784e-04, reg: 2.707126e-04, (17 / 100)
val_acc: 0.509000, lr: 9.752279e-04, reg: 2.850865e-03, (18 / 100)
val_acc: 0.500000, lr: 2.412048e-04, reg: 4.997821e-04, (19 / 100)
val_acc: 0.466000, lr: 1.319314e-04, reg: 1.189915e-02, (20 / 100)
val_acc: 0.516000, lr: 8.039527e-04, reg: 1.528291e-02, (21 / 100)
```

53% - relatively good for a 2-layer neural net with 50 hidden neurons.

23



## Now run finer search...

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-5, 5)
    lr = 10**uniform(-3, -6)
```

adjust range

```
max_count = 100
for count in xrange(max_count):
    reg = 10**uniform(-4, 0)
    lr = 10**uniform(-3, -4)
```

val_acc: 0.527000, lr: 5.340517e-04, reg: 4.097824e-01, (0 / 100)
val_acc: 0.492000, lr: 2.279484e-04, reg: 9.991345e-04, (1 / 100)
val_acc: 0.512000, lr: 8.680827e-04, reg: 1.349727e-02, (2 / 100)
val_acc: 0.461000, lr: 1.028377e-04, reg: 1.220193e-02, (3 / 100)
val_acc: 0.460000, lr: 1.113730e-04, reg: 5.244309e-02, (4 / 100)
val_acc: 0.498000, lr: 9.477776e-04, reg: 2.001293e-03, (5 / 100)
val_acc: 0.469000, lr: 1.484369e-04, reg: 4.328313e-01, (6 / 100)
val_acc: 0.522000, lr: 5.586261e-04, reg: 2.312685e-04, (7 / 100)
val_acc: 0.530000, lr: 5.808183e-04, reg: 8.259964e-02, (8 / 100)
val_acc: 0.489000, lr: 1.979168e-04, reg: 1.010889e-04, (9 / 100)
val_acc: 0.490000, lr: 2.036031e-04, reg: 2.406271e-03, (10 / 100)
val_acc: 0.475000, lr: 2.021162e-04, reg: 2.287807e-01, (11 / 100)
val_acc: 0.460000, lr: 1.135527e-04, reg: 3.905040e-02, (12 / 100)
val_acc: 0.515000, lr: 6.947668e-04, reg: 1.562808e-02, (13 / 100)
val_acc: 0.531000, lr: 9.471549e-04, reg: 1.433895e-03, (14 / 100)
val_acc: 0.509000, lr: 3.140888e-04, reg: 2.857518e-01, (15 / 100)
val_acc: 0.514000, lr: 6.438349e-04, reg: 3.033781e-01, (16 / 100)
val_acc: 0.502000, lr: 3.921784e-04, reg: 2.707126e-04, (17 / 100)
val_acc: 0.509000, lr: 9.752279e-04, reg: 2.850865e-03, (18 / 100)
val_acc: 0.500000, lr: 2.412048e-04, reg: 4.997821e-04, (19 / 100)
val_acc: 0.466000, lr: 1.319314e-04, reg: 1.189915e-02, (20 / 100)
val_acc: 0.516000, lr: 8.039527e-04, reg: 1.528291e-02, (21 / 100)

53% - relatively good for a 2-layer neural net with 50 hidden neurons.

But this best cross-validation result is worrying. Why?

24



## Summary

- Activation Functions (use ReLU)
- Data Preprocessing (images: subtract mean, divide std)
- Weight Initialization (use Xavier/He init)
- Batch Normalization (use)
- Babysitting the training process

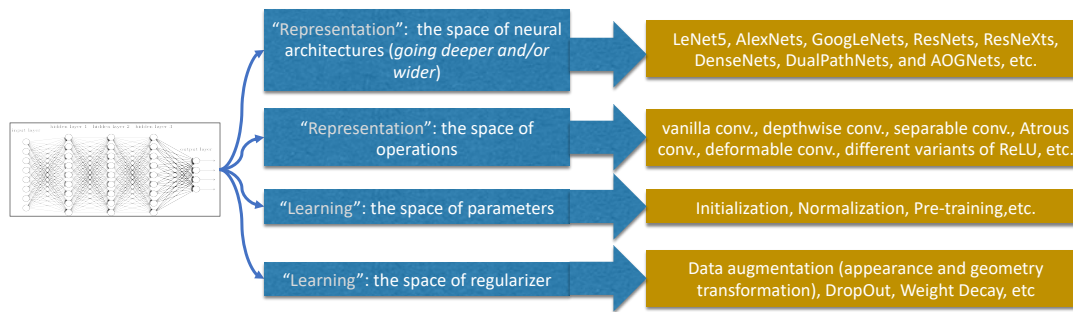
25





## [Optional] Bonus Points for Creativity and Novelty

- Exploring new designs in, but not limited to, three aspects as follows.



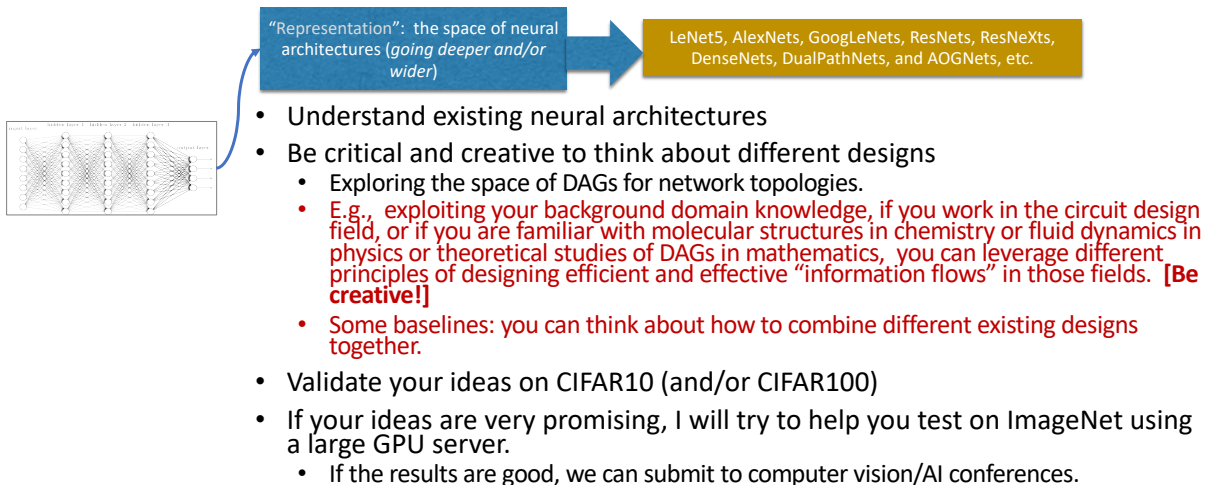
26

26



## Bonus Points for Creativity and Novelty

- Potential Topic 1: New design of network architectures



27

27





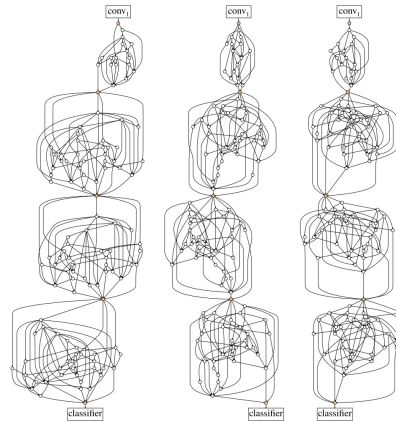
## Bonus Points for Creativity and Novelty

- Potential Topic 1: New design of network architectures
- Examples to inspire you

### Exploring Randomly Wired Neural Networks for Image Recognition

Saining Xie Alexander Kirillov Ross Girshick Kaiming He  
Facebook AI Research (FAIR)

Google and Read the paper



28

28



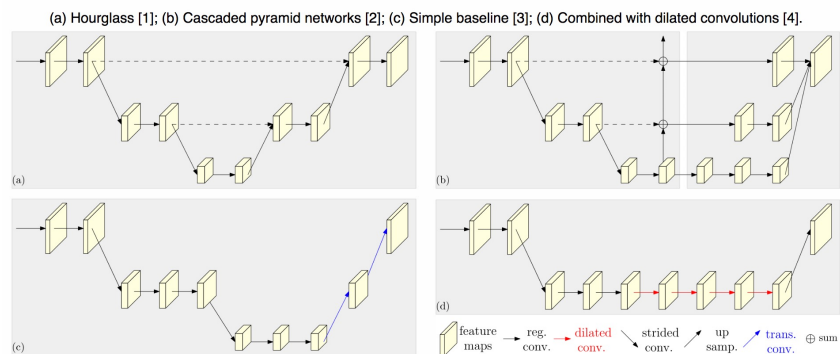
## Bonus Points for Creativity and Novelty

- Potential Topic 1: New design of network architectures
- Examples to inspire you

### Deep High-Resolution Representation Learning for Human Pose Estimation

Ke Sun Bin Xiao Dong Liu [Jingdong Wang](#)

Google and Read the paper



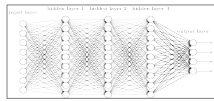
29

29



## Bonus Points for Creativity and Novelty

### Potential Topic 2: New design of node operations



"Representation": the space of operations

vanilla conv., depthwise conv., separable conv., Atrous conv., deformable conv., different variants of ReLU, etc.

- Understand existing operations
- Be critical and creative to think about different designs
  - Consider existing network architectures such as [ResNets](#) or [DenseNets](#) (PyTorch has default implementation), and how to extend their operations (e.g., they use the Bottleneck operation)
  - Exploring the combinations of different primitive operations using some micro-topologies similar in the spirit to network-in-network (<https://arxiv.org/abs/1312.4400>).
  - Examples (understand how they improved the existing networks):
    - SENet: <https://arxiv.org/abs/1709.01507>
    - Non-local networks: <https://arxiv.org/abs/1711.07971> <https://arxiv.org/abs/1811.11721>
    - Image Transformer: <https://arxiv.org/abs/1802.05751>
    - HetConv <https://arxiv.org/pdf/1903.04120.pdf>
  - Exploiting your background domain knowledge on methods of how to explore spatial and channel contextual information, or attention mechanism for operation design. **[Be creative!]**
- Validate your ideas on CIFAR10 (and/or CIFAR100)
- If your ideas are very promising, I will try to help you test on ImageNet or other big datasets using a large GPU server.
  - If the results are good, we can submit to computer vision/AI conferences.

30

30



## Bonus Points for Creativity and Novelty

- Potential Topic 2: New design of node operations
- Examples to inspire you

### Res2Net: A New Multi-scale Backbone Architecture

Shang-Hua Gao\*, Ming-Ming Cheng\*, Kai Zhao, Xin-Yu Zhang, Ming-Hsuan Yang, and Philip Torr

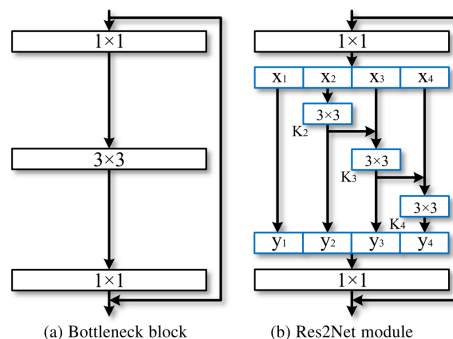


Fig. 2: Comparison between the bottleneck block and the proposed Res2Net module (the scale dimension  $s = 4$ ).

31

31



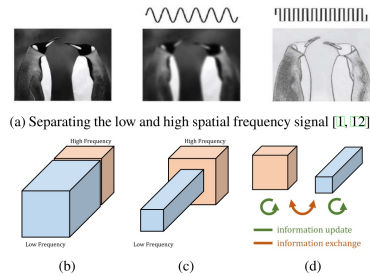
## Bonus Points for Creativity and Novelty

- Potential Topic 2: New design of node operations
- Examples to inspire you

### Drop an Octave: Reducing Spatial Redundancy in Convolutional Neural Networks with Octave Convolution

Yunpeng Chen<sup>1,2</sup>, Haoqi Fang<sup>1</sup>, Bing Xu<sup>1</sup>, Zhicheng Yan<sup>1</sup>, Yannis Kalantidis<sup>1</sup>,  
 Marcus Rohrbach<sup>1</sup>, Shuicheng Yan<sup>1,2</sup>, Jiashi Feng<sup>2</sup>

<sup>1</sup>Facebook AI, <sup>2</sup>National University of Singapore, <sup>3</sup>Qihoo 360 AI Institute



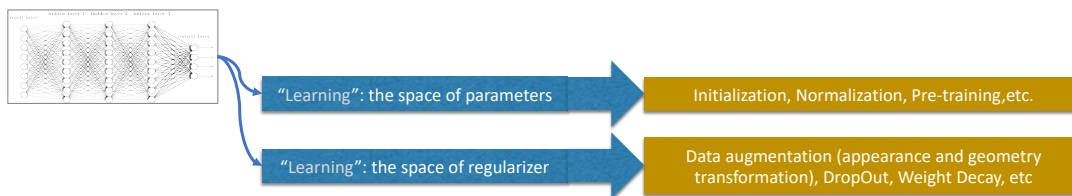
32

32



## Bonus Points for Creativity and Novelty

- Potential Topic 3: New design of learning



- Understand existing initialization and normalization methods
- Be critical and creative to think about different designs
  - Consider existing network architectures such as [ResNets](#) or [DenseNets](#) (PyTorch has default implementation), and how to train them with different initialization/normalization methods
  - Examples (understand how they improved the vanilla random initialization and the vanilla BatchNorm methods):
    - <https://arxiv.org/abs/1502.01852>
    - Fixup: <https://arxiv.org/abs/1901.09321v1>
    - BatchNorm, InstanceNorm, LayerNorm and GroupNorm: <https://arxiv.org/abs/1803.08494> (and related references therein)
    - How does BatchNorm work: <https://arxiv.org/abs/1805.11604>
    - SwitchableNorm: <https://github.com/switchablenorms/Switchable-Normalization>
    - Attentive Norm: <https://arxiv.org/abs/1908.01259>
    - Feature response norm: <https://arxiv.org/pdf/1911.09737.pdf>
- Validate your ideas on CIFAR10 (and/or CIFAR100)
- If your ideas very promising, I will try to help you test on ImageNet or other big datasets using a large GPU server.
  - If the results are good, we can submit to computer vision/AI conferences.

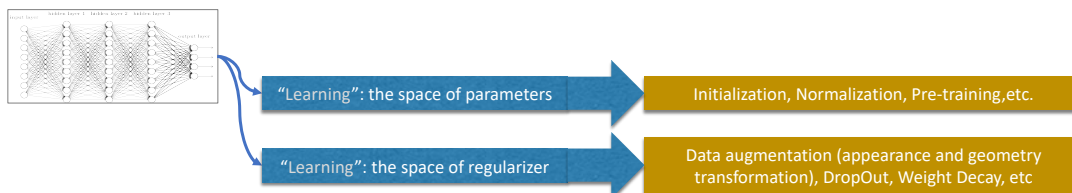
33

33



## Bonus Points for Creativity and Novelty

### • Potential Topic 3: New design of learning



- Understand existing data augmentation "tricks"
- Be critical and creative to think about different designs
  - Consider existing network architectures such as [ResNets](#) or [DenseNets](#) (PyTorch has default implementation), and how to train them with different initialization/normalization methods
  - Examples (understand how they augment the training data):
    - Cutout: <https://arxiv.org/abs/1708.04552>
    - Mixup: <https://arxiv.org/abs/1710.09412>
    - Bag of tricks: <https://arxiv.org/abs/1812.01187> (and references for the tricks therein)
    - Grid Mask: <https://arxiv.org/abs/2001.04086>
  - Any transformation functions that do not change the task semantics are applicable. **[Be creative!]**
- Validate your ideas on CIFAR10 (and/or CIFAR100)
- If your ideas very promising, I will try to help you test on ImageNet or other big datasets using a large GPU server.
  - If the results are good, we can submit to computer vision/AI conferences.

34